

2025



**Survey Assessing Risks from
Artificial Intelligence (SARA) 2025:
Australian Public Attitudes
Toward AI Risks and Governance**

November 2025

aigovernance.org.au

SARA-25

Table of contents

Executive Summary	2
Findings in Full	4
Public Expectations of AI Governance	4
Public Perceptions of AI Risk	10
Public Risk Tolerance	15
Our Recommendations	23
Methodology	25
Survey Design and Implementation	25
Statistical Approach	25
Limitations	27
Methodological Considerations	27
Implications of Limitations	28
References	29
Appendices	31
Appendix A: Survey Questions	31
Appendix B: Full-Label Plots	36
Appendix C: Statistical Details	38
Appendix D: Sample Demographics	39
Appendix E: Demographic Heterogeneity Analysis	40
Appendix F: Robustness Checks	41
Appendix G: Comparison Between Risk Tolerance and Forecast Risk Levels	42
Appendix H: Testing Robustness of Risk Tolerance and Scope Sensitivity	46
Appendix I: Risk Tolerance Among Consistency-Validated Respondents	49

Executive Summary

While AI adoption is argued to be essential for economic competitiveness¹, there is public resistance driven by deep mistrust and perceptions of inadequate regulation². Our [survey](#) of 933 Australians finds the public expects AI to be as safe as commercial aviation - at least **4,000x** safer than current risk estimates. They want the government to better manage AI risks, and that many risk controls would increase their trust in AI.

Key Findings

- **Australians desire regulation:** 74% worry the government [won't regulate AI enough](#), and 83% believe [regulation is lagging behind technological progress](#).
- **Public trust in AI developers is extremely low:** Less than one in four Australians (23%) [trust technology companies](#) to ensure AI safety.
- **Risk management outweighs innovation:** When forced to choose, 72% want the government to [prioritise managing AI risks](#) over driving innovation.
- **Safety expectations as high as for aircraft:** 94% expect AI systems to [meet or exceed the safety standards of commercial aviation](#)—around 4,000 to 30,000 times safer than expert risk estimates.
- **Australians reject catastrophic risks:** Even a 1% [chance of a global catastrophe](#) is considered unacceptable by most respondents.
- **Patience for safety is widespread:** A majority would support [delaying advanced AI development](#) by up to 50 years if it reduced catastrophic risk from 5% to 0.5%.
- **Australians want robust safety measures:** More than 89% [would trust AI more](#) if there were mandatory safety testing, independent audits, or an Australian AI Safety Institute.
- **Australians favour global limits on dangerous AI:** 57% support an international treaty [banning the development of "smarter-than-human" AI](#).
- **Demand for transparency and media coverage is strong:** 80–85% want [more reporting](#) on AI's societal effects and on how government is managing AI regulation.

Recommendations

- **Adopt a precautionary approach to the most dangerous AI:** Given strong support for patience and global cooperation, control development of high-risk systems [until achieving credible safety evidence and public consensus](#).
- **Strengthen AI regulation and oversight:** Establish enforceable safety standards, mandatory testing, and transparent reporting to close the gap between public expectations and current governance.
- **Create an independent Australian AI Safety Institute:** Coordinate safety research, conduct audits, and advise government—building public trust while aligning with international best practices.

Method

We recruited 1,068 Australian adults through an online panel between August 14-30, 2025, and analysed the 933 participants (87%) who passed two attention checks. We used multilevel regression and post-stratification (MRP)—the current standard for converting non-probability samples into population estimates—to weight responses by age, gender, education, and state against 2021 Census distributions.

Declarations of Conflict of Interest

The authors have no conflicts to declare.

Citation

Noetel, M., Saeri, A.K., Graham, J.G. (2025). *Survey Assessing Risks from Artificial Intelligence (SARA) 2025: Australian Public Attitudes Toward AI Risks and Governance*. The University of Queensland.
<https://aigovernance.org.au/survey/2025>

Funding Declaration

This project was funded by the lead author.

Findings in Full

Public Expectations of AI Governance


Low Trust Inhibits AI Adoption

Governments around the world argue that artificial intelligence is critical for economic growth and national security⁴. Yet Australia faces a fundamental challenge: public resistance threatens to undermine AI adoption as countries around the world pull ahead². In this report, we first explore why Australians struggle to trust AI and what regulations might increase that trust.

First, we explored their personal reasons for not using AI. When Australians avoid AI, they say it's because they have privacy concerns (57%), because they prefer doing things without AI (39%), or because they don't trust technology companies (32%) or like tech culture (31%).

Australians cite varied reasons for not adopting AI

"What are the main reasons you don't use AI tools more often?"



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

Figure 1: Barriers preventing Australians from using AI tools

This aligns with global findings where 57% of respondents agree AI poses significant threats to privacy⁵, reflecting what Stanford HAI researchers call "societal-level privacy risks that existing regulatory frameworks are not designed to address"⁶.

Despite low rates of AI literacy², Australians do not report knowledge as a primary barrier^a. This suggests that the public might be unlikely to engage with education around AI unless the core trust issues are resolved.

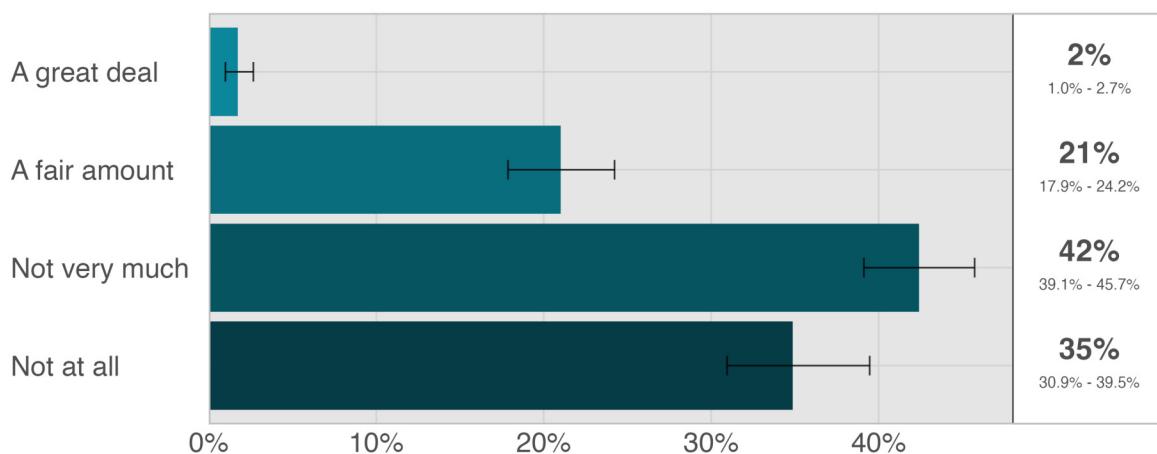
^aOnly 19% cite lack of understanding

Australians Don't Trust Tech-Company Self-Regulation

Currently, those developing or deploying AI in Australia are regulated by general laws designed for traditional technologies¹ rather than a dedicated AI act, like that of the European Union⁷. The Government has proposed voluntary guardrails that hope to align with international standards⁸, but companies would be left to self-enforce these guardrails. The problem is: Australians don't trust the organisations developing and deploying AI.

Australians do not trust tech companies to ensure AI is safe

*"To what extent do you trust tech companies
to ensure the AI they develop is safe?"*



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

Figure 2: Australian trust in tech companies to ensure AI safety

When asked to what extent they trust tech companies to ensure the AI they develop is safe, **77%** of Australians said they trust them "not at all" or "not very much." Only **2%** said they trust tech companies "a great deal."

Compared with an October 2023 international survey funded by the UK government⁹, this places us among the least trusting of surveyed countries. Japan showed similar distrust of technology company self-regulation, which was substantially below Canada, USA, France, Italy, Singapore, the UK, and South Korea.

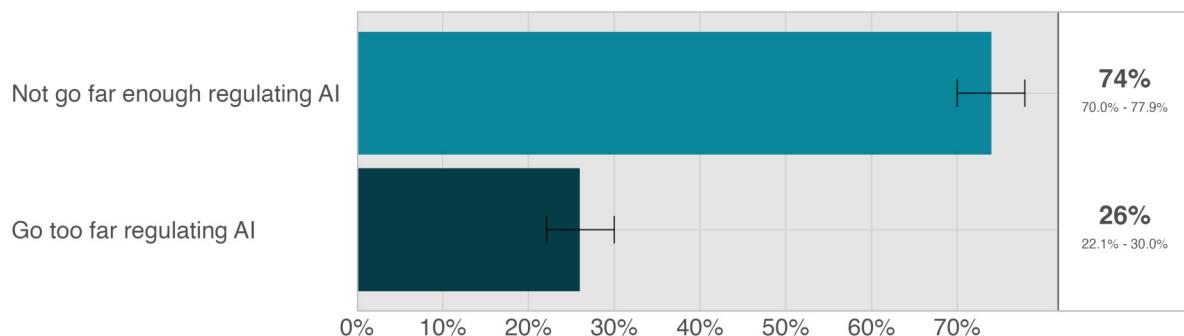
A 48,000-person study across 47 countries found systematically low confidence in commercial organisations developing AI in the public interest². That study also showed Australians had particularly low trust in those organisations (5th lowest of 47 countries). Theoretical frameworks suggest that this trust deficit stems from gaps between ethical principles and actual practices¹⁰, or from governance failures¹¹.

Australians Worry Their Government Won't Go Far Enough

Our data suggest the public are more concerned about the government failing to put in enough regulation, than regulating too much. When forced to choose, **74%** said they worry the government won't regulate AI enough, while only **26%** fear over-regulation.

Australians worry government will not regulate AI enough

*"Thinking about the use of artificial intelligence (AI) in Australia,
are you more concerned that the Australian government will go too far
regulating its use or not go far enough regulating its use?"*



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

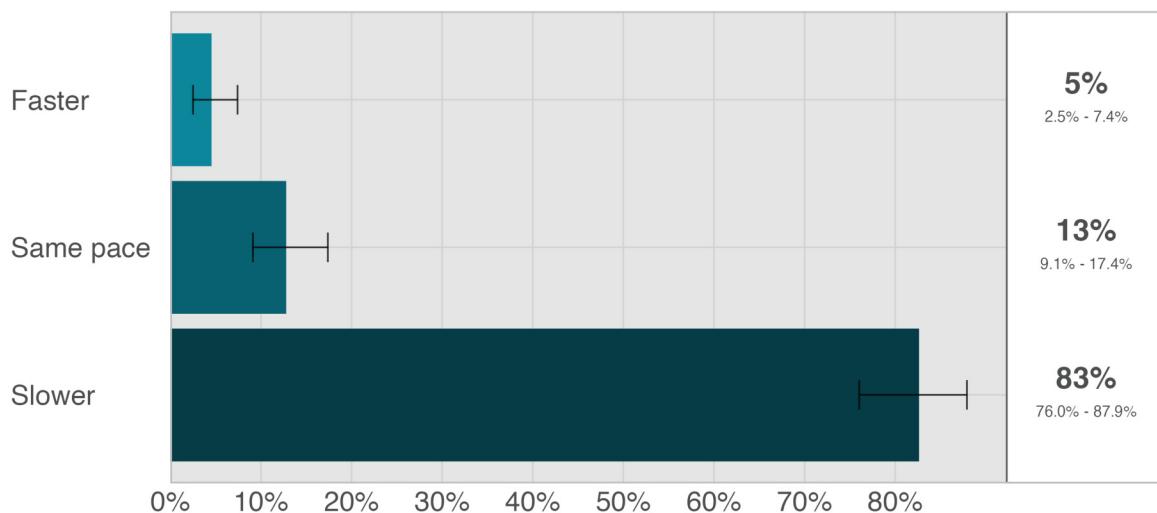
Figure 3: Australian concerns about AI regulation pace

Australians Think Regulation is Lagging Behind the Technology

The public also judge the regulatory response to be slow. Only 5% feel that regulation is developing ahead of the technology, and only 13% feel that regulation is keeping pace. 83% of Australians think regulation is falling behind technological innovation. This mirrors UK polling¹² that shows only 11% of people think AI is keeping pace with AI technologies.

Australians believe AI regulation is lagging behind technology

“Is regulation developing faster or slower than AI technology?”



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

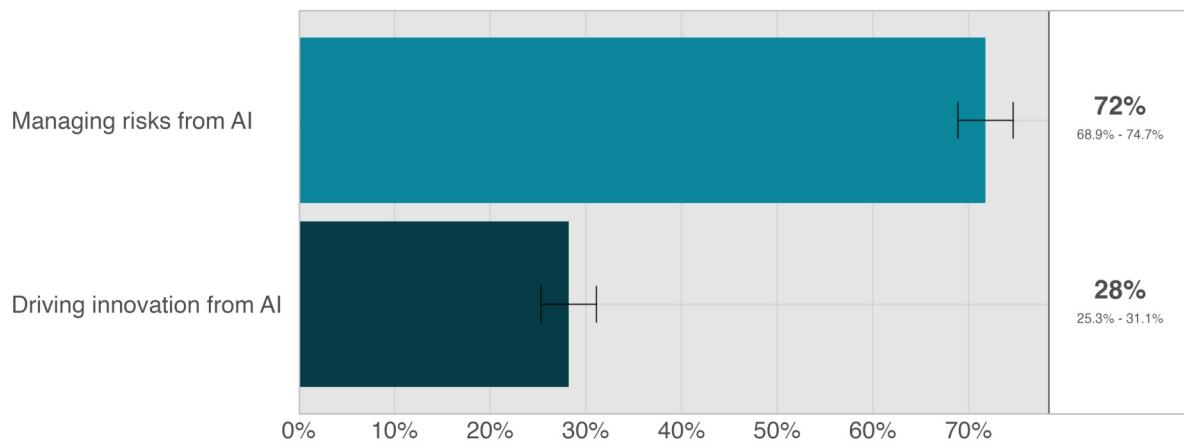
Figure 4: Perceived pace of AI regulation relative to technology development

Australians Expect the Government to Prioritise Addressing Risks

The Federal Government¹ describes tensions between promoting innovation and managing risks. When asked whether the government should focus on managing risks or driving innovation, **72%** of Australians said the government's priority should be risk management. This preference aligns with the broader pattern: the public wants government to protect the public as AI is developed and deployed.

Australians prefer government prioritise risk management over innovation

"If forced to prioritise, the Australian government should focus on:"



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

Figure 5: Public preferences for government AI priorities

Public Perceptions of AI Risk


Most Risks Are Seen As Priorities

Australians did not have clear priorities for which risks they want the government to support. We drew risks from the international AI safety report^{13 b}. At the end, we also asked whether they felt the government should prioritise the risks on the public's behalf, or prioritise economic adaptation following AI deployment.

Ten different risks showed strong agreement (>87%), from privacy (96%) to cyber attacks (93%) and loss of control (89%). The public were less concerned about equitable global access to AI (53%) and energy use (73%). They were relatively unsupportive of experts focusing on adaptation (65% agreement; i.e., 'helping the economy adapt to the change') and expert prioritisation (76%).

Australians think many AI risks are priorities

"In managing risks from AI, I think the government should focus on..."



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

Figure 6: AI risks Australians believe government should prioritise

^bWe used Claude Opus 4 to summarise those risks at a Year 7 reading age (full items in Figure 19), and shuffled the list for each participant.

Australians Are Increasingly Uncertain About AI's Net Impact

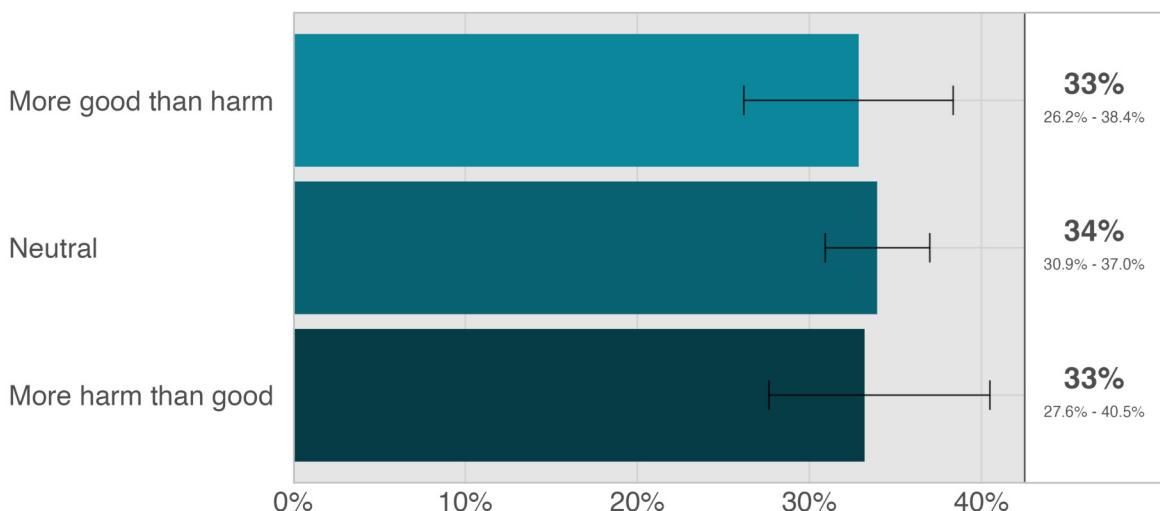
When asked whether AI will do more good or harm overall, Australians in 2025 are essentially split three ways:

- 33.3% believe AI will do more harm than good,
- 33.9% think harms will balance benefits, and
- 32.7% think it will do more good than harm.

This represents a significant shift from 2024, when the public was more polarized. In our 2024 study¹⁴, only 20% were neutral. This year, we saw a large increase in uncertainty (+13.9 percentage points), with corresponding decreases in both pessimistic (-9.7pp) and optimistic (-4.3pp) views.

Australians divided on whether AI will do more harm or good

“Will AI do more good or harm overall?”



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

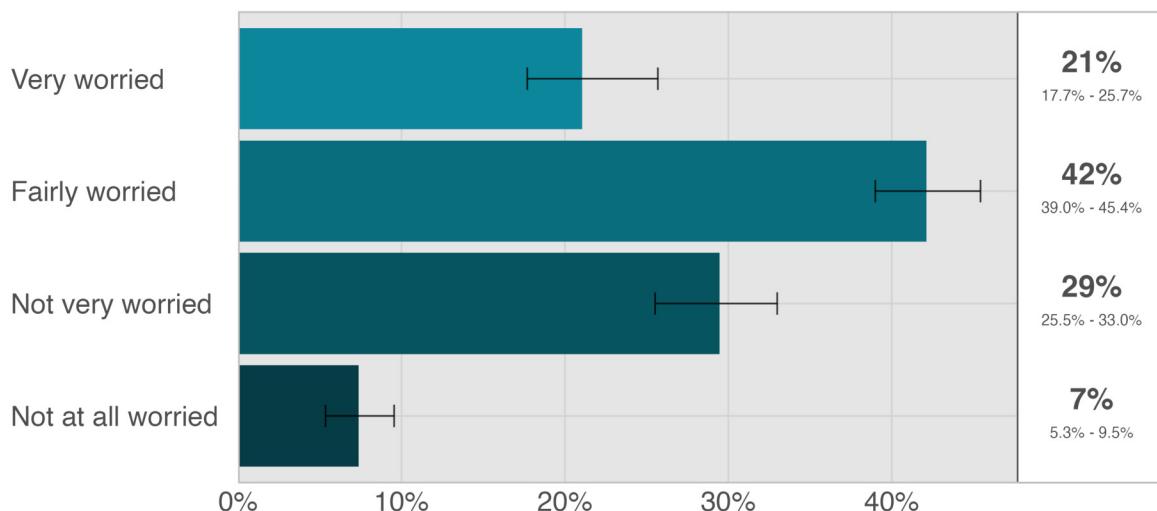
Figure 7: Growing uncertainty: Australians shift from polarized to ambivalent views on AI's impact

Australians Worry About Losing Jobs to AI

As a salient specific harm, we looked at concern about job loss. A recent Harvard study showed that companies that adopt AI have already reduced hiring for junior roles by 7.7% since 2023¹⁵. When asked how worried they are that AI will lead to widespread unemployment, **63%** of Australians said they are "fairly worried" or "very worried" about job loss from AI. This mirrors US data showing 56% of Americans are 'extremely or very concerned' about AI leading to job loss¹⁶.

Australians worry about AI-driven job loss

"How worried are you about AI leading to job loss?"



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

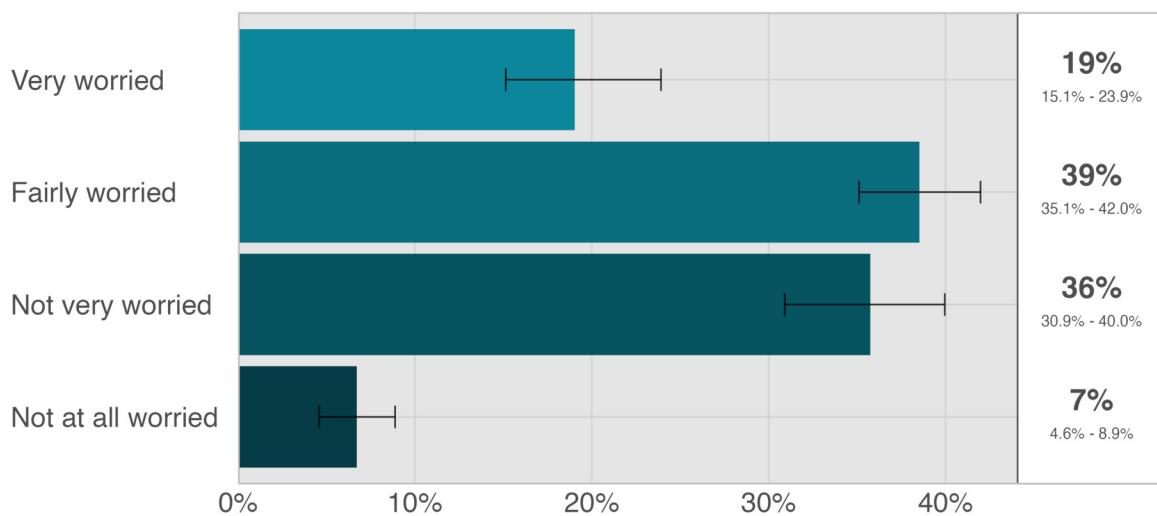
Figure 8: Level of worry about unemployment from AI

Australians Worry About Losing Control of AI

As a more extreme risk, international AI experts worry¹³ that humanity could lose control of powerful frontier AI systems. When asked how worried they are that humans will lose control of AI, **58%** said they are "fairly worried" or "very worried."

Australians worry about losing control of AI

"How worried, if at all, are you that humans will lose control of AI?"



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

Figure 9: Level of worry about humans losing control of AI

These risks have different probabilities and severities^c, yet elicit similar levels of concern. This could reflect several factors:

1. Both widespread unemployment and loss-of-control represent fundamental disruptions to social order that people find deeply concerning regardless of probability.
2. The abstract nature of 'loss-of-control'^d may lead people to underestimate its severity compared to the concrete, familiar threat of job loss, or
3. Australians may be implicitly weighting probability and severity, judging that near-certain economic disruption warrants similar concern as possible catastrophic outcomes.

^cUnemployment is common and present; loss-of-control is speculative and unprecedented. Unemployment is debilitating but, in Australia, not life-threatening. In the worst instance, loss-of-control could mean the end of humanity¹⁷

^dOur previous research suggests concrete images increases perceived risks¹⁸. Without concrete explanations, it might not be obvious that 'loss of control' leads to any deaths. People could think it means an AI system doing something we don't want for some short time, like purchasing the wrong size shoes.

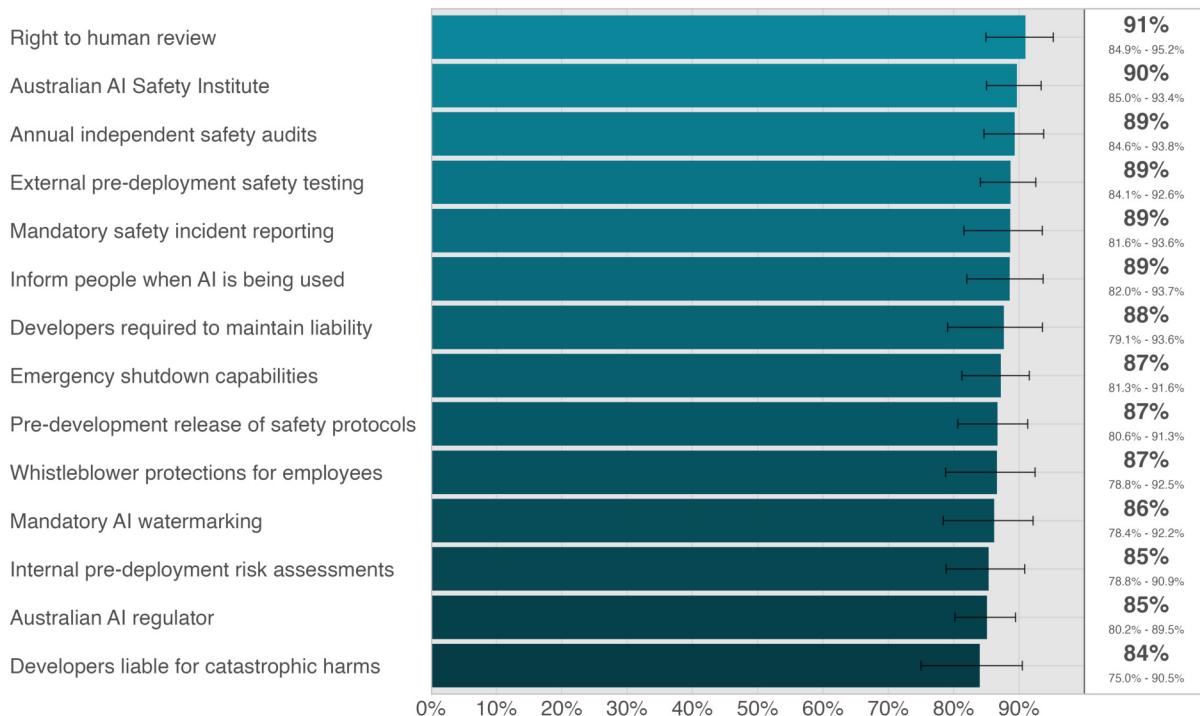
Most Mitigations Would Reportedly Increase Trust

We collected mitigations from various government proposals¹⁹⁻²² to see if they would increase the public's trust in AI.^e

14 of the 15 mitigations we tested would make most Australians more likely to trust AI, with the right to human review (91%) and an Australian AI Safety Institute (90%) topping the list. Independent safety audits also attract overwhelming support (89%).

Australians endorsed many measures to increase trust in AI

"I would be more likely to trust AI if..."



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

Figure 10: Measures that would increase Australian trust in AI

Our estimates of population attitudes has some uncertainty (see confidence intervals on the right). As a result, it is hard to definitively rank the proposed measures using the data above. The strong support for most measures appears consistent with other polling. For example, pollsters have found between 60%²³ and 80%²⁴ of US voters support mandatory testing, safety protocols, and liability for frontier AI model developers.

These indicate that the government likely has many levers through which it could increase trust and safety.

^eWe used Claude Opus 4 to describe these mitigations in Year 10 English (full labels in appendices as [list](#) and [figure](#)) and shuffled the list for each participant.

Public Risk Tolerance

Safety Standards and Risk Tolerance

In many areas of public safety (e.g., civil engineering, power generation, aviation), regulators set acceptable safety standards and certify engineers to design systems that meet those standards²⁵. In this model, the public and their elected representatives might not be consulted on the specific safeguards put in place. Instead, if they are consulted, regulators might instead mostly consider their risk tolerance. For example, they might not have input in the design of a dam, but could rightly expect the dam to have a very low chance of bursting. The regulator is expected to make that risk transparent so the public can make informed decisions (e.g., to move, or to request a lower risk tolerance).

For AI safety, a similar approach would be understand the levels of societal risk the public would be willing to tolerate, then let regulators put in place the safeguards that bring risk down to those levels. As noted above, it is still important to understand which safeguards would increase the public's trust, if trust is important for increasing adoption. However, they are not technical experts who understand which mitigations actually reduce risk. For example, they might say that 'watermarking' AI-generated content would increase their trust without understanding how easy it is to circumvent watermarks^f.

If the goal is both real and perceived trust, it is important to understand stakeholder preferences and expert judgements about which mitigations reduce risk to tolerable levels. In this section, we outline the public's expectations of those levels.

^f"We prove that, under well-specified and natural assumptions, strong watermarking is impossible to achieve. This holds even in the private detection algorithm setting, where the watermark insertion and detection algorithms share a secret key, unknown to the attacker."²⁶
p. 1

Australians Expect AI To Be As Safe As Commercial Aviation

We asked whether AI should have standards more or less strict than aviation. 94% of Australians want advanced AI to meet aviation-grade safety or better, including 58% who prefer standards stricter than airlines and just 6% willing to accept looser rules.


This is far safer than current estimates from forecasting and AI experts^{27,28} who put the risk of 8 billion deaths between 1% and 10%. Leading AI researchers, and the CEOs of [frontier AI companies](#), estimate catastrophic risks from advanced AI between 2% and 25%.

Even the most optimistic credible assessment (superforecasters) shows AI risk at ~**4,000x** aviation levels. The median domain expert assessment suggests ~**30,000 times** higher risk than aviation (see Appendix G).

There is a gulf between expectations and reality: The public expects AI to match the safety of an industry that hasn't had a major Australian commercial airline fatality, while experts, forecasters, and AI CEOs assess it as carrying extinction-level risks.

Australians expect AI safety standards at least as strict as aviation

“Different technologies have different safety standards. Compared to commercial airlines (less than 1 in 10,000,000 risk of death per flight), advanced AI systems should be...”



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

Figure 11: Desired AI safety standards compared to aviation

We explored directly asking about risk tolerances for different sized risks (see [Appendix H](#)). There were trade-offs between making the question simple and making small probabilities easy to understand. These probabilities are difficult to validly elicit from the public, but findings generally support claims above: people expect safety standards akin to commercial aviation.

Even Radical Benefits Are Not Worth Catastrophic Risks

We also tested whether major benefits could justify catastrophic risks. Each respondent was randomly assigned a risk level (0.01% to 20%) and asked: "An advanced AI system has a [X]% chance of causing a global catastrophe (over 1 million deaths) but a 60% chance of solving climate change and extending human lifespan by 20 years. Is this trade-off acceptable?"

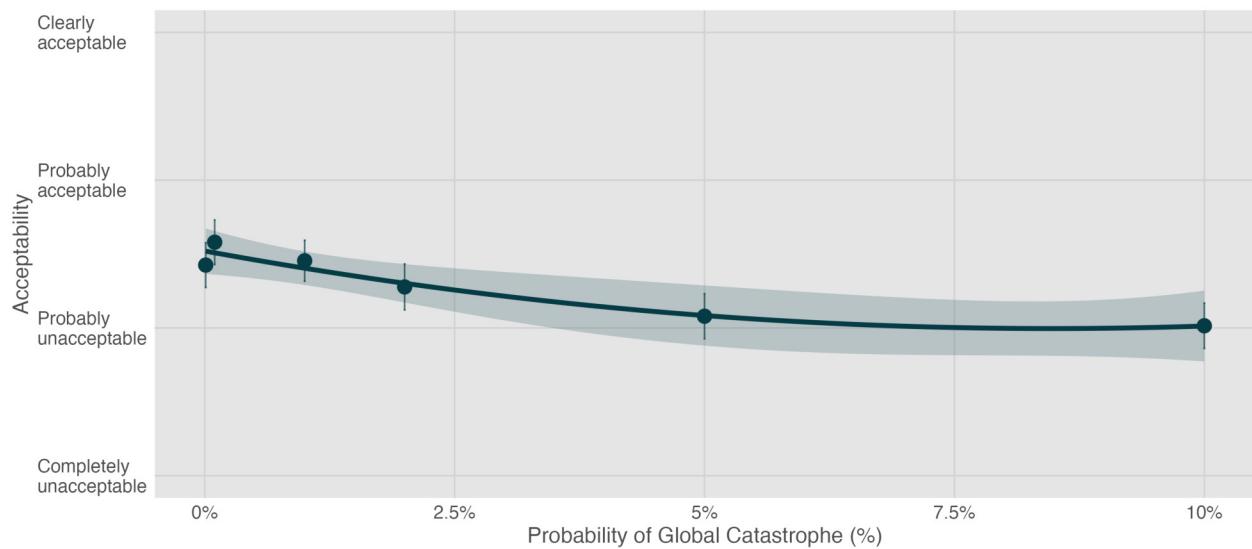
As seen below, even small probabilities of catastrophic outcomes are contentious for the public. No level of catastrophic risk is clearly acceptable.

Australians accept almost no catastrophic risk even for dramatic AI benefits

An advanced AI system has a ___% chance of causing a global catastrophe (over 1 million deaths)

but a 60% chance of solving climate change and extending human lifespan by 20 years.

This trade-off is:



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

Figure 12: Acceptability of catastrophic AI risks for major benefits

At 0.01% risk (1 in 10,000), **46% reject** the trade-off. At 1% risk, **51% find it unacceptable**. This indicates the public are divided on whether AI that could 'solve climate change and extend lifespan' is worth even a 1 in 1,000 risk of a catastrophe.


Australians Show Patience for Safety

Question Design: We measured time preferences for AI safety by randomizing delay periods (1 to 50 years). Respondents were asked: "Suppose we could reduce the risk of AI catastrophe from 5% to 0.5% by delaying advanced AI development by [X] years. Would this delay be worthwhile?"

There was strong support for delays across all timeframes tested (57% to 87%). Most Australians (80%) would support a 10-year delay, and even **50-year delays receive majority support** (57%), with only 8% saying such delays were not worthwhile. Australians prioritise safety over speed, willing to wait generations for safer AI rather than accept 5% catastrophic risk.

Australians willing to wait for AI benefits

If delaying advanced AI deployment by ___ years would reduce catastrophic risk from 5% to 0.5%, this delay would be:



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

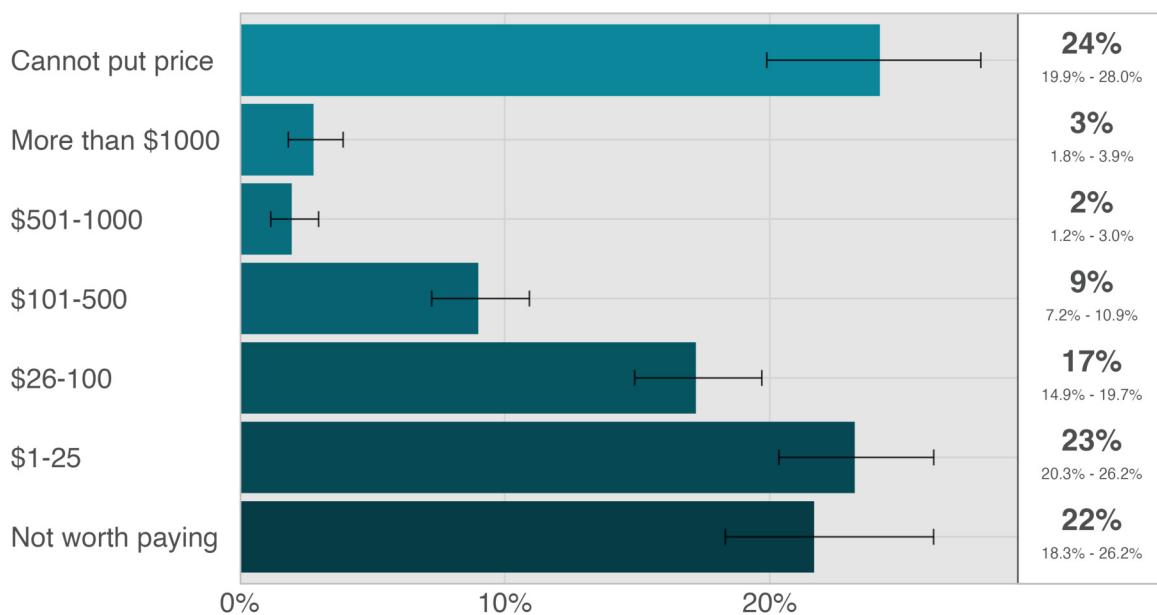
Figure 13: Willingness to delay AI deployment for safety improvements

Australians Are Divided on Paying for Safety

Our data show 22% of Australians say the safety improvements are **not worth paying for**, while a similar share (24%) say they **cannot put a price** on preventing catastrophe. About 40% would contribute a modest annual amount (up to \$100), and only 14% are ready to pay more than \$100 each year. Overall, roughly half of the public (54%) is willing to pay something to cut the risk from 1% to 0.01%, but contributions cluster at the lower end (23% would pay just \$1–25, and 17% would pay \$26–100).

Australians show mixed willingness to pay for AI risk reduction

“What is the maximum you would personally pay annually (through taxes/prices) to reduce AI catastrophic risk (over 1 million deaths) from 1% to 0.01% over the next 30 years?”



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).


Figure 14: Willingness to pay annually to reduce catastrophic AI risk from 1% to 0.01%. Bars show MRP-adjusted population estimates with 95% credible intervals. N=933 respondents who passed attention checks.

Support for Global Ban on Artificial Superintelligence

The International AI Safety Report²⁹ suggests that substantial risk of catastrophic outcomes stems from difficulty controlling smarter-than-human systems. While this is a goal of many frontier AI companies (e.g., ChatGPT's OpenAI), **57%** of Australians support an international treaty banning the development of 'smarter-than-human' artificial intelligence (a.k.a., 'artificial superintelligence').

Australians support a ban on artificial superintelligence

"To what extent would you support or oppose the introduction of an international treaty to ban any 'smarter-than-human' artificial intelligence (AI)?"



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).


Figure 15: Support for international treaty banning superintelligent AI

Public Request for More Coverage

Beyond specific risks, Australians want ongoing public discourse about AI's societal impacts and AI regulation. 79% of Australians want more media coverage of AI's societal impacts and 85% want more reporting on how government is regulating AI; barely 21% and 15% say they do **not** need additional coverage.

Australians want more media coverage of AI's societal impact

*"I want to hear more from the media about
how AI is going to affect society"*

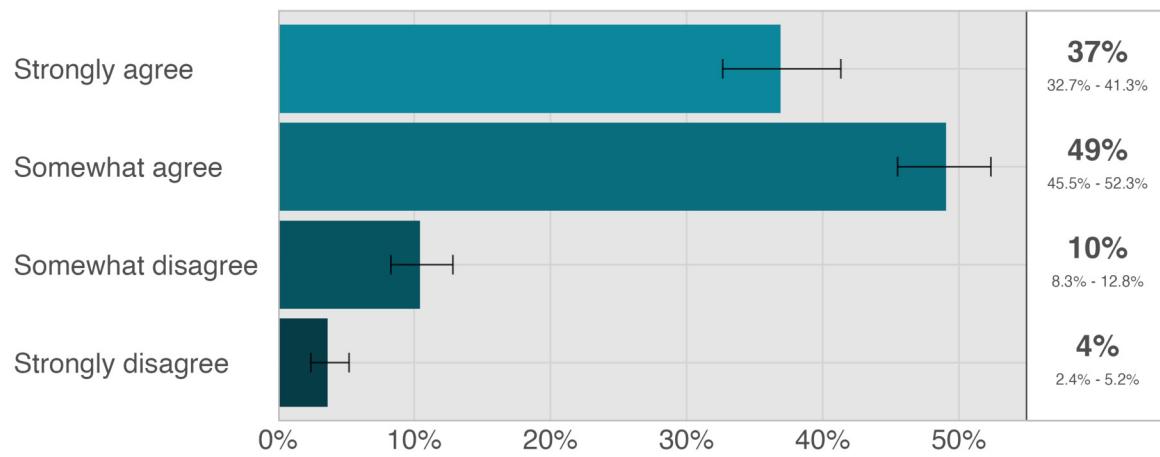


Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

Figure 16: Desire for media coverage of AI's societal impact

Australians want more media coverage of AI regulation

*"I want to hear more from the media about
how the government is regulating AI"*



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

Figure 17: Desire for media coverage of AI regulation

Our Recommendations

Technology-Specific Regulation Is Needed to Meet Public Safety Expectations

Bodies like the Australian Productivity Commission argue¹ that technology-specific regulations should be 'a last resort' because badly designed or heavy handed rules would slow adoption. Poorly coordinated regulations are cited as adoption headwinds for global firms³⁰. In our view, the Productivity Commission correctly identifies that Australia's existing legal frameworks⁹ provide substantial scaffolding for AI governance¹. Companies that steward data responsibly, ensure algorithmic fairness, and maintain transparency may build trust incrementally.

However, our findings suggest the public's perceptions of slow or weak regulation might be inhibiting trust and adoption. For example:

- They fear that AI currently has [insufficient privacy protections](#), that the regulation is [lagging behind the technology](#), and think the government is [more likely to under- than over-regulate](#).
- They think the government should [prioritise risk management](#) over accelerating adoption, and [see many risks as priorities](#). For example, 63% worry about [losing jobs](#) and 58% [losing control of AI itself](#).
- Australians expect stringent standards on AI, akin to the [commercial aviation industry](#). [Expert assessments](#) suggest risks over the coming decades are at least 4,000x higher than this expectation.
- If it reduces risks, Australians are [willing to wait for advanced AI](#), and about [half seem willing to pay](#).
- They [want to hear more](#) about how AI is affecting society and what is being done to regulate it.

Technology-Specific Regulation Would Likely Improve Safety and Trust

AI creates unprecedented risks²⁹. For example, leading AI companies plan to build 'agentic' general models "that can autonomously plan, act, and pursue its own goals across almost all tasks that humans can perform"³¹. Capabilities like these pose catastrophic risks²⁹. The forecast risks from AI are much higher than the public is willing to tolerate. As a result, governments around the world are implementing a range of technology-specific safeguards⁷.

Many such safeguards focus on catastrophic risks from frontier AI development²⁰. They focus on these risks due to the size of the threat and because more comprehensive regulation has been more controversial³². The regulations aim to increase safety without imposing unnecessary standards on smaller companies and low-risk uses of AI³³. Our survey data shows that many mitigations would reportedly increase the public's trust. For example, they say they would be more likely to trust AI if the government implemented SB53's key provisions:

- Require incident reporting
- Require published safety protocols
- Protect whistleblowers from retaliation

⁹e.g., privacy laws, consumer protections, anti-discrimination legislation

They also felt it would increase trust if the government had an AI Safety Institute to understand risk and work with industry, as they have in the [UK](#), [USA](#) and [Korea](#). These mitigations might be useful first steps that monitor and reduce risk while bringing Australia into line with emerging standards.

Instead of aiming to only meet emerging standards, Australia may need to look ahead and lead in some areas, particularly given how quickly AI is developing. As noted above, the public are more concerned about the government being too slow and doing too little. Our previous work¹⁴ showed 8 in 10 want Australia to lead in international governance of AI. Regulators could therefore consider the other popular mitigations (e.g., independent safety testing, emergency shutdown capabilities, developer liability for harms) argued to reduce catastrophic risks. There are emerging dialogues about which additional mitigations would be most effective amid rapid progress^{13,29}.

There are also emerging dialogues around the need to pursue agentic general AI models at all³⁴. 25,000 signatories have called for "a prohibition on the development of superintelligence, not lifted before there is broad scientific consensus that it will be done safely and controllably, and strong public buy-in." We found majority support (57%) for an even more strict "international treaty to ban any 'smarter-than-human' AI". Given such prohibitions likely require international coordination³⁵, Australia should have a continuing conversation about whether or not its representatives should support such a prohibition.

Overall, with [so many risks to prioritise](#), and deep distrust of the technology², countries like Australia have a difficult job. Significant work is required to reduce risks to levels expected by the public. Still, our findings point to many democratically popular approaches for making AI both trusted and trustworthy.

Methodology

Survey Design and Implementation

Sample Recruitment

The survey was conducted with **1069 Australian adults** between **August 14-30, 2025**, recruited through Prolific to achieve demographic diversity. We used Prolific because comparison studies have suggested it provides some of the best data quality³⁶.

Quality Control

- Two attention check questions embedded in survey
- **933 participants (87.3%)** passed both checks
- These 933 responses form our analysis sample
- Validated Australian residential postcodes only
- Median completion time analysis to identify 'speeders'

Randomisation

Key experimental manipulations:

- Module presentation order
- Casualty numbers (10 to 8 billion)
- Risk percentages (0.01% to 20%)
- Delay periods (1 to 50 years)

Statistical Approach

Multilevel Regression and Post-stratification (MRP)

MRP is the state-of-the-art method for generating population-representative estimates from non-probability samples^{37,38}.

How MRP Works:

1. **Model individual responses** using demographic and geographic predictors
2. **Generate predictions** for all demographic-geographic combinations
3. **Weight predictions** using Census population data

4. **Aggregate** to population-level estimates with uncertainty

Advantages:

- Corrects for any sampling bias
- Provides uncertainty quantification
- Validated in election, health, and public opinion contexts
- More accurate than simple weighting

Technical Specifications

- **Software:** R with `brms` package
- **Model:** Bayesian multilevel regression
- **Iterations:** 2000 per chain, 1000 warmup
- **Chains:** 4 parallel chains
- **Convergence:** R-hat < 1.01, ESS > 400
- **Post-stratification:** 2021 Australian Census data

Limitations

Methodological Considerations

While this study employs state-of-the-art methods, several limitations warrant careful consideration when interpreting results.

Selection Bias

Our online panel recruitment methodology likely oversamples individuals comfortable with technology:

- **Tech-savvy bias:** Respondents willing to complete online surveys may be more familiar with AI than the general population
- **Underestimated concerns:** If tech-comfortable individuals express this level of concern (74% worry about under-regulation), actual population concerns may be even higher
- **Limited reach:** Digitally excluded populations—often older, rural, or socioeconomically disadvantaged—are underrepresented despite MRP adjustments

This selection bias suggests our findings may represent a *lower bound* on public concerns about AI risks.

Attention Check Validity

The 12.7% attention check failure rate raises important questions:

- **Different population segment?** Attention check failures might represent a distinct group—perhaps those more sceptical of surveys or less engaged with technical topics
- **Conservative analysis:** By excluding these responses, we may have removed legitimate voices, particularly from less educated or engaged segments

Future research should explore whether attention check “failures” represent a meaningful population subset with different AI attitudes.

Question Framing Effects

Several aspects of our question design may have influenced responses:

Anchoring effects:

- Specifying “global catastrophe (over 1 million deaths)” provides a concrete anchor that may shape risk perception

- Comparing to known risks (aviation, nuclear) frames AI in terms of established dangerous technologies

Hypothetical scenarios:

- Asking about international treaties or 50-year delays involves speculation about unfamiliar concepts
- Public may agree with precautionary measures without fully understanding implications

Implications of Limitations

Despite these limitations, the findings provide valuable insights:

1. **Conservative estimates:** Selection bias may mean we *underestimate* public concerns
2. **Robust patterns:** The consistency of findings across measures suggests real underlying attitudes
3. **Policy relevance:** Even if effects are partially driven by framing, they reflect how public will respond to real policy debates
4. **Action imperative:** Limitations don't negate the clear trust deficit and demand for governance

References

1. Productivity Comission. *Harnessing Data and Digital Technology*. <https://www.pc.gov.au/inquiries/current/data-digital/interim/data-digital-interim.pdf> (2025).
2. Gillespie, N., Lockey, S. & Curtis, C. *Trust, Attitudes and Use of Artificial Intelligence: A Global Study 2025*. (2025).
3. Office of Science and Technology Policy. *Winning the Race: America's AI Action Plan*. <https://www.ai.gov/action-plan> (2025).
4. OECD/BCG/INSEAD. *The Adoption of Artificial Intelligence in Firms: New Evidence for Policymaking*. https://www.oecd.org/en/publications/the-adoption-of-artificial-intelligence-in-firms_f9gef33c3-en.html (2025) doi:10.1787/f9gef33c3-en.
5. International Association of Privacy Professionals. *Consumer Perspectives of Privacy and Artificial Intelligence*. (2024).
6. King, T. & Meinhardt, N. *Rethinking Privacy in the AI Era: Policy Provocations for a Data-Centric World*. <https://hai.stanford.edu/white-paper-rethinking-privacy-ai-era> (2024).
7. European Union. Regulation (EU) 2024/1689 Laying down Harmonised Rules on Artificial Intelligence (Artificial Intelligence Act). (2024).
8. Resources, D. of I. S. and. *Voluntary AI Safety Standard | Department of Industry Science and Resources*. <https://www.industry.gov.au/node/93303> (2025).
9. Department for Science, Innovation and Technology and Centre for Data Ethics and Innovation. *International survey of public opinion on AI safety*. GOV.UK.
10. Floridi, L. et al. *AI4People—An Ethical Framework for a Good AI Society: Opportunities, Risks, Principles, and Recommendations*. *Minds and Machines* **28**, 689–707 (2018).
11. Winfield, A. F. T. & Jiroka, M. *Ethical Governance Is Essential to Building Trust in Robotics and Artificial Intelligence Systems*. *Philosophical Transactions of the Royal Society A* **376**, 20180085 (2018).
12. Ipsos. *Public Attitudes to AI*. <https://www.ipsos.com/sites/default/files/ct/publication/documents/2023-11/Public%20Perceptions%20of%20AI%20charts.pdf> (2023).
13. Bengio, Y. et al. *International AI Safety Report*. <https://www.gov.uk/government/publications/international-ai-safety-report-2025> (2025).
14. Saeri, A., Noetel, M. & Graham, J. Survey Assessing Risks from Artificial Intelligence (Technical Report). (2024) doi:10.2139/ssrn.4750953.
15. Lichtinger, G. & Hosseini Maasoum, S. M. Generative AI as Seniority-Biased Technological Change: Evidence from U.S. Résumé and Job Posting Data. (2025) doi:10.2139/ssrn.5425555.
16. McClain, C., Kennedy, B., Gottfried, J., Anderson, M. & Pasquini, G. *How the U.S. Public and AI Experts View Artificial Intelligence*. Pew Research Center (2025).
17. Daniel Kokotajlo, Scott Alexander, Thomas Larsen, Eli Lifland & Romeo Dean. *AI 2027*. AI 2027 (2025).
18. Ward, T., Saeri, A. & Noetel, M. Turn a Blind AI: The Impact of Compassion Fade and the Identifiable Victim Effect on AI Risk Concerns. (2024) doi:10.2139/ssrn.4839839.
19. Department of Industry, Science and Resources. *Safe and Responsible AI in Australia: Proposals Paper for Introducing Mandatory Guardrails for AI in High-Risk Settings*. https://storage.googleapis.com/converlens-au-industry/p/prj2f6fo2ebfe6a8190c7bdc/page/proposals_paper_for_introducing_mandatory_guardrails_for_ai_in_high_risk_settings.pdf (2024).

20. Senator Wiener. [Artificial intelligence models: Large developers](#). (2025).
21. Micah Lasher, Rebecca Seawright, Amy Paulin & Yudelka Tapia. [NY State Assembly Bill 2025-A6453A](#). (2025).
22. Department of Industry Science and Resources. [The Seoul Declaration by countries attending the AI Seoul Summit, 21-22 May 2024](#). *Department of Industry, Science and Resources* (2024).
23. Artificial Intelligence Policy Institute. [2.5x More Voters Support SB1047 than Oppose it in Collaborative Poll Between Anti-Bill Economist and AIPI - AI Policy Institute](#). (2024).
24. Artificial Intelligence Policy Institute. [Poll: Californians Support Strong Version of SB1047, Disagree With Anthropic's Proposed Changes - AI Policy Institute](#). (2024).
25. International Standards Organisation. [ISO 31000:2018](#). (2018).
26. Zhang, H. et al. Watermarks in the Sand: Impossibility of Strong Watermarking for Generative Models. (2025) doi:[10.48550/arXiv.2311.04378](https://doi.org/10.48550/arXiv.2311.04378).
27. Karger, E. et al. [Forecasting Existential Risks: Evidence from a Long-Run Forecasting Tournament](#). (2023).
28. Grace, K. et al. [Thousands of AI Authors on the Future of AI](#). arXiv preprint arXiv:2401.02843 (2024).
29. Bengio, Y. et al. [Managing AI Risks in an Era of Rapid Progress](#). *Science* **384**, 842–845 (2024).
30. Renieris, E. M., Kiron, D. & Mills, S. [A Fragmented Landscape Is No Excuse for Global Companies Serious About Responsible AI](#). *MIT Sloan Management Review* (2024).
31. Bengio, Y. et al. Superintelligent Agents Pose Catastrophic Risks: Can Scientist AI Offer a Safer Path? (2025) doi:[10.48550/arXiv.2502.15657](https://doi.org/10.48550/arXiv.2502.15657).
32. [California Just Passed the First U.S. Frontier AI Law. Here's What It Does](#). *Carnegie Endowment for International Peace*.
33. Shen, O. & Lim, J. [Landmark laws for AI transparency](#). *United States Studies Centre* (2025).
34. Future of Life Institute. [Statement on Superintelligence. Statement on Superintelligence](#).
35. United Nations. [Governing AI for Humanity: Final Report](#). (United Nations, New York, NY, 2024).
36. Peer, E., Rothschild, D., Gordon, A., Evernden, Z. & Damer, E. [Data quality of platforms and panels for online behavioral research](#). *Behavior Research Methods* **54**, 1643–1662 (2022).
37. Downes, M. et al. [Multilevel Regression and Poststratification: A Modeling Approach to Estimating Population Quantities From Highly Selected Survey Samples](#). *American Journal of Epidemiology* **187**, 1780–1790 (2018).
38. Gelman, A., Lax, J., Phillips, J., Gabry, J. & Trangucci, R. [Using Multilevel Regression and Poststratification to Estimate Dynamic Public Opinion](#). (2018).
39. Said, H. [IATA Annual Safety Report - 2024](#).
40. Grace, K., Stein-Perlman, Z., Weinstein-Raun, B. & Salvatier, J. [2022 Expert Survey on Progress in AI](#). <https://aiimpacts.org/2022-expert-survey-on-progress-in-ai/> (2022).
41. Department of Planning. [HIPAP 4: Risk Criteria for Land Use Safety Planning](#). (2011).
42. Lundgren, R. E. & McMakin, A. H. Visual Representations of Risks. in *Risk Communication* 159–191 (John Wiley & Sons, Ltd, 2013). doi:[10.1002/9781118645734.ch14](https://doi.org/10.1002/9781118645734.ch14).
43. Visschers, V. H. M., Meertens, R. M., Passchier, W. W. F. & De Vries, N. N. K. [Probability Information in Risk Communication: A Review of the Research Literature](#). *Risk Analysis* **29**, 267–287 (2009).

Appendices

Appendix A: Survey Questions

Core Attitude Questions

AI Good vs Harm "Overall, do you think artificial intelligence (AI) will do more good or more harm?"

- More good than harm
- Neutral
- More harm than good

Trust in Tech Companies "To what extent do you trust tech companies to ensure the AI they develop is safe?"

- A great deal
- A fair amount
- Not very much
- Not at all
- Don't know

Worry About Control "How worried, if at all, are you that humans will lose control of AI?"

- Very worried
- Fairly worried
- Not very worried
- Not at all worried
- Don't know

Worry About Job Loss "When it comes to artificial intelligence (AI), how worried are you about AI leading to job loss?"

- Very worried
- Fairly worried
- Not very worried
- Not at all worried
- Don't know

International Treaty on Advanced AI "To what extent would you support or oppose the introduction of an international treaty to ban any 'smarter-than-human' artificial intelligence (AI)?"

- Strongly support

- Somewhat support
- Somewhat oppose
- Strongly oppose
- Don't know

Regulation Pace "Regulation of AI in Australia is developing _____ than development in AI technologies."

- Faster
- Same Pace
- Slower
- Don't know

Regulation Concern "Thinking about the use of artificial intelligence (AI) in Australia, are you more concerned that the Australian government will go too far regulating its use or not go far enough regulating its use?"

- Go too far regulating its use
- Not go far enough regulating its use
- Not sure

Media Coverage of AI "I want to hear more from the media about how AI is going to affect society."

- Strongly agree
- Somewhat agree
- Somewhat disagree
- Strongly disagree
- Not sure

Media Coverage of AI Regulation "I want to hear more from the media about how the government is regulating AI."

- Strongly agree
- Somewhat agree
- Somewhat disagree
- Strongly disagree
- Not sure

Attention Check "This one is not about AI. I have recently visited the moon."

- Strongly agree
- Somewhat agree
- Somewhat disagree
- Strongly disagree
- Not sure

Barriers to AI Use "What are the main reasons you don't use AI tools more often? (Select all that apply)"

- I don't trust the companies that make them
- I'm concerned about my data privacy
- I think AI tools are unsafe or risky
- I don't understand how to use them

- I'm worried about meeting my legal or professional obligations
- I don't like the culture around AI and tech companies
- They're not relevant to my work or interests
- I prefer doing things without AI assistance
- They're too expensive
- I'm concerned they might replace human jobs
- Other

Government Focus "If forced to prioritise, the Australian government should focus on:"

- Managing risks from AI
- Driving innovation from AI

Trust-Building Measures

For each measure, participants were asked: "I would be more likely to trust AI if..."

- there was an Australian AI Safety Institute that helps understand advanced AI, brief the Government on its risks, and work with industry on deploying it safely
- there was an Australian AI regulator who sets and enforces laws and safety standards
- technology companies were required to submit their most powerful AI models to a regulator for safety testing before being deployed
- developers of frontier AI models were required to create and publish detailed safety and security protocols before building the system
- large AI developers were required to undergo annual independent audits of their safety compliance
- AI developers were required to report safety incidents (like model escapes, unauthorised access, or autonomous behaviour) to authorities within 72 hours
- employees were protected from retaliation when reporting AI safety concerns, even if their employer is technically following the law
- there were rules preventing AI developers from making contracts that remove their liability for harms
- developers of large AI models were liable for catastrophic harms caused by their models
- AI models had emergency shutdown capabilities such that model developers can disable their model in case of an emergency
- organisations were required to conduct risk assessments before deploying AI systems that could impact human rights, health and safety, or have significant societal effects
- organisations were required to inform people when AI was being used to make decisions about them (such as in hiring, lending, or government services)
- AI-generated content (text, images, audio, video) was required to be labeled or watermarked so people knew it was created by AI
- people had the right to challenge decisions made by AI systems and have their complaints reviewed by a human

Risk Priorities

"In managing risks from AI, I think the government should focus on..."

- stopping bad actors from using AI to create fake videos or images that hurt people (like scams or fake nude photos)
- stopping people from using AI-made fake content to trick the public
- stopping hackers from using AI to attack computers, businesses, and important services like power grids
- stopping AI from making it easier to create dangerous weapons (like biological or chemical weapons)

- making sure AI systems work properly and don't make mistakes or give false information
- stopping AI systems from being unfair or treating some groups of people worse than others
- preventing situations where AI systems could get out of human control
- dealing with how AI affects people's jobs and employment
- making sure all countries have fair access to AI technology (not just rich countries)
- preventing problems that happen when only a few big companies control all the AI technology
- reducing the environmental damage from AI's energy use
- protecting people's privacy from AI that collects personal data or watches what people do
- dealing with problems when AI uses other people's work without permission
- I just want experts to focus on preventing whatever has the biggest risk
- I just want experts to focus on helping the economy adapt to the change

Risk Tolerance Questions

Casualty Risk (with randomised casualty number N) "Some experts think that advanced AI could potentially cause accidents where people die. What is the HIGHEST probability you would tolerate of AI causing [N] casualties each year?"

- 1% or higher (10,000+ in 100,000 chance) Similar to risk of a global pandemic
- 0.1% (100 in 100,000 chance) Similar to risk of death tolerated for workplace accidents
- 0.01% (10 in 100,000 chance) Similar to risk of death tolerated from dam failure
- 0.001% (1 in 100,000 chance) Similar to risk of death tolerated from nuclear power
- 0.0001% (0.1 in 100,000 chance) Similar to risk of death tolerated from commercial planes
- 0.00001% or less (<0.01 in 100,000 chance) Similar to risk of extinction-level asteroid
- Zero risk only is acceptable
- This is too speculative to answer meaningfully

Risk-Benefit Trade-off (with randomised risk percentage X) "An advanced AI system has a |X|% chance of causing a global catastrophe (over 1 million deaths) but a 60% chance of solving climate change and extending human lifespan by 20 years. This trade-off is:"

- Completely unacceptable
- Probably unacceptable
- Probably acceptable
- Clearly acceptable
- Cannot say

AI versus Existing Technologies "Different technologies have different safety standards. Compared to commercial airlines (less than 1 in 10,000,000 risk of death per flight), advanced AI systems should be:"

- Much stricter safety standards
- Somewhat stricter standards
- Same safety standards
- Somewhat less strict
- Much less strict
- Cannot compare these technologies

Safety Standards Attention Check "Various technologies require safety testing. To ensure you're reading carefully, please select 'Somewhat stricter than nuclear power' for this question about safety standards."

- Much stricter than nuclear power
- Somewhat stricter than nuclear power

- Same as nuclear power
- Somewhat less strict than nuclear power
- Much less strict than nuclear power
- No standards needed

Willingness to Pay for Safety "What is the maximum you would personally pay annually (through taxes/prices) to reduce AI catastrophic risk (over 1 million deaths) from 1% to 0.01% over the next 30 years?"

- \$0 - Not worth paying for
- \$1-25
- \$26-100
- \$101-500
- \$501-1000
- More than \$1000
- Cannot put price on this

Temporal Trade-off "If delaying advanced AI deployment by [Delay] years would reduce catastrophic risk from 5% to 0.5%, this delay would be:"

- Clearly worthwhile
- Probably worthwhile
- Difficult to judge
- Probably not worthwhile
- Clearly not worthwhile


Consistency Validation "Earlier you indicated [Threshold] was your maximum acceptable risk for [Casualty] deaths from AI. To confirm, a technology with a [Threshold] chance of causing [Casualty] deaths is:"

- Far too risky
 - Somewhat too risky
 - About right
 - Could accept somewhat higher risk
 - Could accept much higher risk
-

Appendix B: Full-Label Plots

Australians endorsed many measures to increase trust in AI

"I would be more likely to trust AI if..."

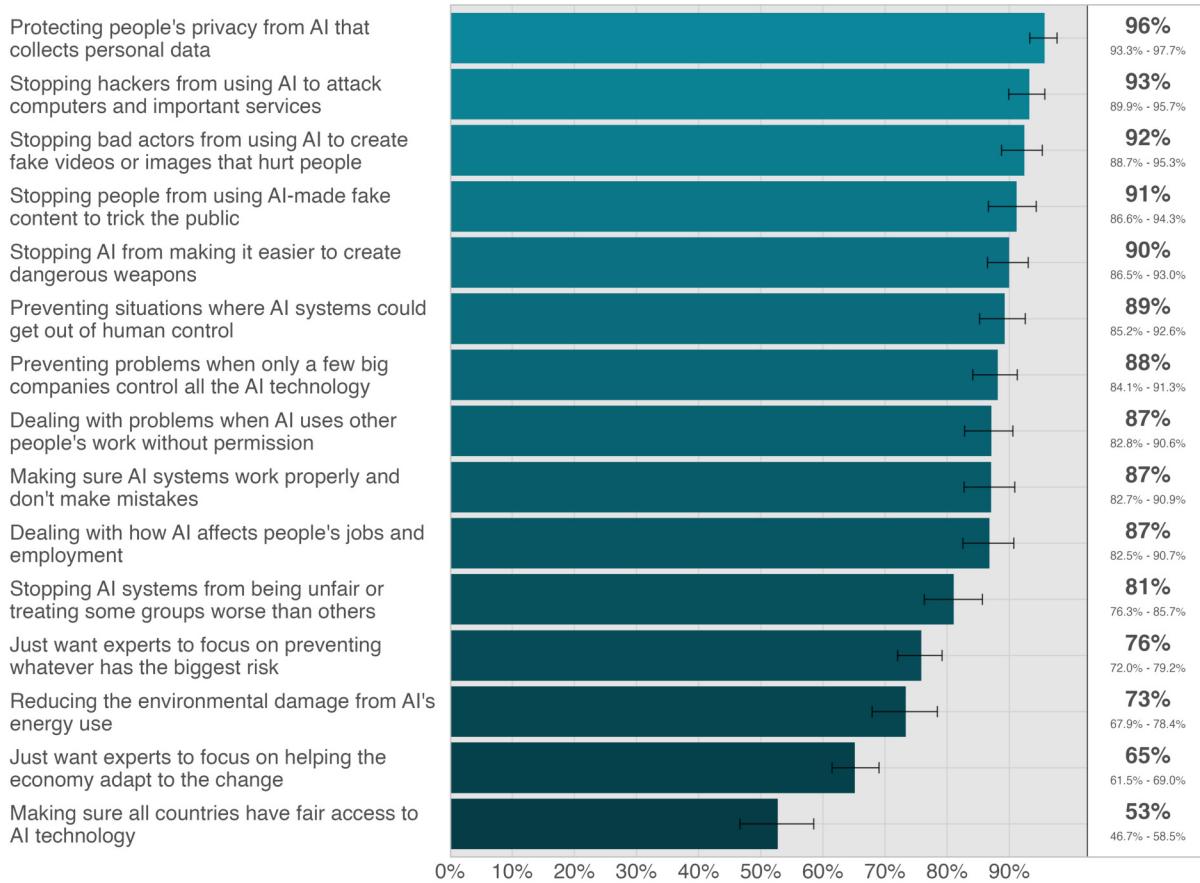


Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

Figure 18: Trust measures with complete question wording

Australians think many AI risks are priorities

"In managing risks from AI, I think the government should focus on..."



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.
Population estimates using Multilevel Regression and Post-stratification (MRP).

Figure 19: Priority risks with complete question wording

Appendix C: Statistical Details

Model Specifications

Binary Outcomes Model:

```
y ~ 1 + (1 | state) + (1 | age_group) + (1 | education) + (1 | gender)
```

Ordinal Outcomes Model:

```
y ~ 1 + (1 | state) + (1 | age_group) + (1 | education) + (1 | gender)  
family = cumulative(link = "logit")
```

Convergence Diagnostics

All models achieved:

- R-hat values < 1.01 for all parameters
 - Effective sample size > 400 for all parameters
 - No divergent transitions
 - Visual inspection of trace plots showed good mixing
-

Appendix D: Sample Demographics

Sample Characteristics

Table 1: Sample demographics compared to 2021 Australian Census

Characteristic	Sample (n=932)	Census 2021	Difference
Age			
18-34	49.1%	32.0%	+17.1pp
35-54	39.3%	35.0%	+4.3pp
55+	11.6%	33.0%	-21.4pp
Gender			
Woman	49.6%	50.0%	-0.4pp
Man	49.5%	50.0%	-0.5pp
Other/Prefer not to say	1.0%	-	-
Education			
No university degree	32.1%	65.0%	-32.9pp
University degree	67.9%	35.0%	+32.9pp
State			
VIC	31.5%	26.0%	+5.5pp
NSW	28.0%	32.0%	-4.0pp
QLD	19.0%	20.0%	-1.0pp
SA	8.5%	7.0%	+1.5pp
WA	7.2%	10.0%	-2.8pp
TAS/ACT/NT	5.8%	5.0%	+0.8pp

The sample shows reasonable demographic representativeness with slight oversampling of younger, university-educated respondents, typical for online panels. Multi-level regression and post-stratification methods are robust to small deviations from census proportions like these.

Appendix E: Demographic Heterogeneity Analysis

Key Attitudes by Age Group

Table 2: Support for key measures by age group with 95% confidence intervals

Question	Sample Size_18-34	Sample Size_35-54	Sample Size_55+	Estimate [95% CI]_18-34	Estimate [95% CI]_35-54	Estimate [95% CI]_55+
Trust tech companies (% trust)	451	353	107	27.7% [23.6-31.8]	20.4% [16.2-24.6]	19.6% [12.1-27.2]
Worried about loss of control (% worried)	442	357	106	50.9% [46.2-55.6]	56.3% [51.2-61.4]	59.4% [50.1-68.8]
Worried about job loss (% worried)	449	363	107	69.5% [65.2-73.7]	57.3% [52.2-62.4]	51.4% [41.9-60.9]
Support AGI treaty ban (% support)	381	283	93	56.2% [51.2-61.2]	54.1% [48.3-59.9]	58.1% [48.0-68.1]

Key Attitudes by Education Level

Table 3: Support for key measures by education level

Question	n_No university degree	n_University degree	estimate_No university degree	estimate_University degree
Trust tech companies (% trust)	289	622	22.5% [17.7-27.3]	24.6% [21.2-28.0]
Worried about loss of control (% worried)	289	616	58.5% [52.8-64.2]	51.9% [48.0-55.9]
Worried about job loss (% worried)	292	627	62.7% [57.1-68.2]	62.5% [58.7-66.3]
Support AGI treaty ban (% support)	244	513	57.0% [50.8-63.2]	55.0% [50.7-59.3]

Appendix F: Robustness Checks

Attention Check Performance

All main analyses use only participants who passed both attention checks (n=933). This table shows the distribution of attention check performance in the full sample.

Table 4: Response differences by attention check performance

Attention Check Performance	N	% of Total Sample
Failed both	26	2.4%
Passed both	933	87.3%
Passed one	110	10.3%

Robustness Finding: Excluding attention check failures does not materially change results. Main estimates shift by <3 percentage points.

Appendix G: Comparison Between Risk Tolerance and Forecast Risk Levels

Setting the Baseline: Aviation Safety Standards

Commercial aviation achieves:

- **144 deaths globally per year** from commercial flights (IATA 2024 rolling average³⁹)
- Global population: ~8 billion
- **Individual annual risk: 1.8×10^{-8} per person per year** ($144/8,000,000,000$)

This is the safety standard the public expects for AI.

Converting XPT Forecasts to Annual Per-Person Risk

From the XPT data²⁷, we have cumulative risks by 2100 (approximately 75 years from now):

Catastrophic Risk (>10% of humanity dies)

- **Domain experts:** 12% cumulative risk by 2100
- **Superforecasters:** 2.13% cumulative risk by 2100

Extinction Risk (humanity extinct or <5000 survivors)

- **Domain experts:** 3% cumulative risk by 2100
- **Superforecasters:** 0.38% cumulative risk by 2100

Step 1: Annualizing Cumulative Risk

For a cumulative risk P over n years, assume the constant annual risk r is:

$$1 - (1 - r)^n = P \text{ therefore: } r = 1 - (1 - P)^{1/n}$$

Step 2: Catastrophic Risk Calculations

Domain Experts:

- $r = 1 - (1 - 0.12)^{1/75} = 1 - 0.88^{1/75} \approx 0.00170$ annual probability
- If catastrophe occurs, conservatively assuming minimum 10% die
- Individual risk: $0.00170 \times 0.1 = 1.70 \times 10^{-4}$ per person per year

Superforecasters:

- $r = 1 - (1 - 0.0213)^{1/75} = 1 - 0.9787^{1/75} \approx 0.000287$ annual probability
- Individual risk: $0.000287 \times 0.1 = 2.87 \times 10^{-5}$ per person per year

Step 3: Extinction Risk Calculations

Domain Experts:

- $r = 1 - (1 - 0.03)^{(1/75)} = 1 - 0.97^{(1/75)} \approx 0.000408$ annual probability
- Individual risk (100% mortality): 4.08×10^4 per person per year

Superforecasters:

- $r = 1 - (1 - 0.0038)^{(1/75)} = 1 - 0.9962^{(1/75)} \approx 0.0000508$ annual probability
- Individual risk: 5.08×10^5 per person per year

Step 4: Combined Risk Assessment

Extinction events are a subset of catastrophic scenarios. To avoid double counting while still using conservative assumptions (10% fatalities for non-extinction catastrophes, 100% for extinction):

Domain Experts Minimum Annual Per-Person Death Risk:

- Non-extinction catastrophic component: $(0.00170 - 0.000408) \times 0.1 \approx 1.29 \times 10^4$
- Extinction component: 4.08×10^4
- **Combined minimum: $\sim 5.37 \times 10^4$**

Superforecasters Minimum Annual Per-Person Death Risk:

- Non-extinction catastrophic component: $(0.000287 - 0.0000508) \times 0.1 \approx 2.36 \times 10^5$
- Extinction component: 5.08×10^5
- **Combined minimum: $\sim 7.44 \times 10^5$**

Key Comparison

Comparing to aviation safety standard of 1.8×10^8 per person per year:

Primary Comparison

- **Domain Expert Assessment:** $5.37 \times 10^4 \div 1.8 \times 10^8 = \sim 29,833 \text{ times higher}$ than aviation safety
- **Superforecaster Assessment:** $7.44 \times 10^5 \div 1.8 \times 10^8 = \sim 4,133 \text{ times higher}$ than aviation safety

Contextualizing the Gap

To put this in perspective:

- **Commercial Aviation:** Kills ~144 people per year globally (IATA 2024 rolling average)
- **AI (Domain Expert Assessment):** Expected damage value (EDV = fatalities x probability) of ~4.3 million people per year on average ($5.37 \times 10^4 \times 8 \text{ billion}$)
- **AI (Superforecaster Assessment):** Expected damage value (EDV) of ~595,000 people per year on average ($7.44 \times 10^5 \times 8 \text{ billion}$)

Conclusion

The current expert-assessed AI risk is **4,000 to 30,000 times higher** than aviation safety standards that the public expects.

Why the XPT Numbers Are Trustworthy: Tournament Design and Participant Credibility

The Tournament Structure

The **Existential Risk Persuasion Tournament (XPT)**²⁷ was a rigorous, multi-stage forecasting tournament conducted in 2022-2023 with 169 participants who spent months developing and refining their risk assessments. The tournament was explicitly designed to produce the highest-quality possible estimates of existential risks.

Key Credibility Factors

1. Superforecasters: Proven Track Record

The tournament included **88 superforecasters**: individuals with demonstrated exceptional accuracy in predicting geopolitical and economic events. These are people who:

- Have been empirically validated through years of forecasting tournaments
- Beat prediction markets and expert panels on near-term questions
- Were recruited from the top performers in previous tournaments run by Good Judgment Inc.

2. Domain Experts: Deep Technical Knowledge

The tournament included **30 AI domain experts**, including:

- Employees at major AI companies (names withheld for confidentiality)
- Academic researchers specialising in AI safety and capabilities

3. Incentive Structure for Accuracy

Participants were:

- **Financially incentivized** for accurate forecasts (prizes for best performers)
- **Scored on calibration** - not just point estimates but also confidence intervals
- **Rewarded for persuasion** - bonuses for convincing others, incentivizing sharing of best arguments

4. Extensive Deliberation Process

The tournament involved:

- **Multiple stages** over several months
- **Millions of words** of written rationales and arguments
- **Thousands of forecasts** with detailed justifications

- Team formation where forecasters could **debate and challenge each other**

This iterative process typically leads to more accurate assessments than one-off surveys.

5. Transparency in Reasoning

Unlike black-box predictions, the XPT required participants to:

- Provide detailed written rationales for their forecasts
- Explicitly state their assumptions
- Respond to counterarguments
- Show their reasoning chains

This transparency allows us to evaluate the quality of reasoning, not just trust the numbers.

Calibration Against Prior Expert Surveys

The XPT results align with other expert assessments, such as the 2022⁴⁰ and 2024²⁸ researcher survey estimating a 5% median chance of AI causing human extinction.

Methodological Strengths

1. **Proper scoring rules** that mathematically incentivize honest probability assessments
2. **Separation of different risk levels** (catastrophic vs. extinction)
3. **Multiple time horizons** (2030, 2050, 2100) allowing for consistency checks
4. **Cross-domain comparison** (AI risks assessed alongside nuclear, bio, and other risks)

Conservative Aspects of These Estimates

Several factors suggest these estimates might actually be *conservative*:

1. **Selection bias toward sceptics:** Superforecasters are selected for accuracy on near-term, conventional events - not tail risks
 2. **Pre-GPT-4 era:** The tournament concluded before recent AI capabilities advances
-

Appendix H: Testing Robustness of Risk Tolerance and Scope Sensitivity

One goal of this survey was to quantify the public's risk tolerances around AI. As noted above, risk tolerances help regulators decide what precautions are required to meet the public's expectations⁴¹. These tolerances are often presented as Frequency-Consequence graphs, because the public naturally expect a lower risk of bigger threats. For example, the following are the criteria used by the New South Wales government for land use⁴¹.

It says that an 'intolerable' risk of 1 death should be 10^{-3} or 0.001 or 1 in 1,000. Safety planning tries to keep the risk below the intolerable levels. It says no risk of more than 1,000 deaths is 'tolerable'. A 'negligible' risk of 1 death is 10^{-5} or 0.00001 or 1 in 100,000. As this level of safety is usually expensive, so risk management usually aims to keep risks in the ALARP region (as low as reasonably practical).

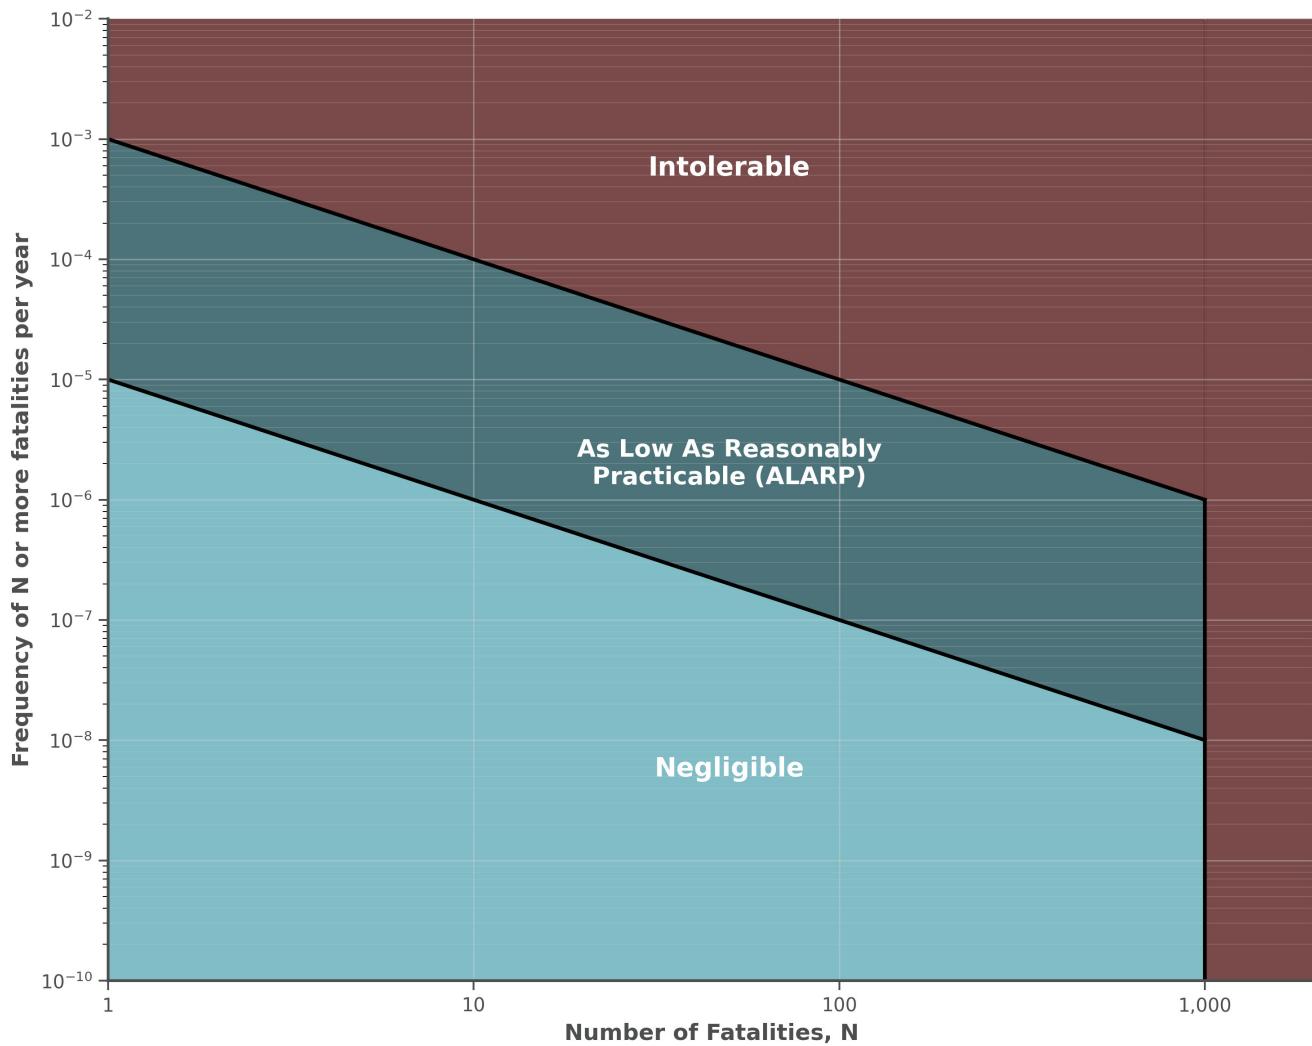


Figure 20: New South Wales Government's Societal Risk Criteria for Land Use Safety Planning


We aimed to explore whether such curves could be drafted for AI.

Question Design: Each respondent was randomly shown a potential casualty number (ranging from 10 people to 8 billion people) and asked: "What is the HIGHEST probability you would tolerate of AI causing (x) casualties each year?" Participants responded via a risk ladder. Risk ladders increase understanding of risks by presenting probabilities in multiple formats, presented vertically in a ladder⁴². Our risk ladder drew on recommendations around risk communication⁴³ by presenting risks as both percentages and natural frequencies with consistent denominators (x in 100,000), anchored to comparable risks (societal risks from dam failure, nuclear power, commercial aviation).

As above, the public expect AI to have the same failure rate as commercial aviation (1 in 1,000,000 chance of death per year). As seen below, Australians were scope insensitive^h, selecting very similar risk tolerances regardless of the number of casualties.

Australians tolerate very low risks of AI accidents (~1 in 1,000,000 / year)

Some experts think that advanced AI could potentially cause accidents where people die. What is the HIGHEST probability you would tolerate of AI causing _ casualties each year?



Sample of 1,069 Australian adults, 933 who passed data quality checks. Fieldwork: August 14-30, 2025.

Population estimates using Multilevel Regression and Post-stratification (MRP).

Figure 21: Maximum acceptable annual risk levels for AI-caused casualties. Each point represents the median response for that casualty level, with error bars showing 95% confidence intervals. N=722 respondents (210 of 933 who passed attention checks selected 'Too speculative to answer'), with casualty numbers randomised.

Note: We included a consistency check which asked the same question with similar wording to confirm if people gave the same judgement. As per [Appendix I](#), 71% of people passed this consistency check, and findings were similar when only including these people.

^h"Scope neglect or scope insensitivity is a cognitive bias that occurs when the valuation of a problem is not valued with a multiplicative relationship to its size."—[Wikipedia](#)

However, we have less confidence in this question due to ambiguities introduced by the anchors on the risk ladder. These anchors were designed to help participants understand the numerical probabilities (e.g., 1 in 100,000). However, the societal anchors people received were the regulated risk tolerances for a **single death** in the [workplace](#), from a [dam](#), from a [nuclear plant](#), or from [flying](#). These were not scaled when the question asked about multiple deaths. For example, someone asked about the risk tolerance of 1 billion deaths was given the same risk anchors (e.g., "0.0001% (0.1 in 100,000 chance) Similar to risk of death tolerated from commercial planes") even though there is not a 0.0001% chance of aircraft causing 1 billion deaths. This introduced ambiguity in how the anchors could be interpreted, particularly for larger numbers.

Nevertheless, we report this item here for completeness as exploratory data. The data generally support findings from [simpler questions](#) showing low risk tolerances from AI.

Appendix I: Risk Tolerance Among Consistency-Validated Respondents

To validate our risk tolerance findings, we re-analysed the randomised experiments using only participants who demonstrated response consistency. The consistency check asked participants to confirm their earlier risk tolerance: "Earlier you indicated [participant selected risk, e.g., 0.01%] was your maximum acceptable risk for [randomly selected number; e.g., 1,000] deaths from AI. To confirm, a technology with a [participant selected risk] chance of causing [casualties] deaths is: Far too risky Somewhat too risky About right Could accept somewhat higher risk Could accept much higher risk" We considered participants as 'inconsistent' when they felt their previous response was 'far too risky' or said they 'could accept much higher risk'.

Consistency Check Results

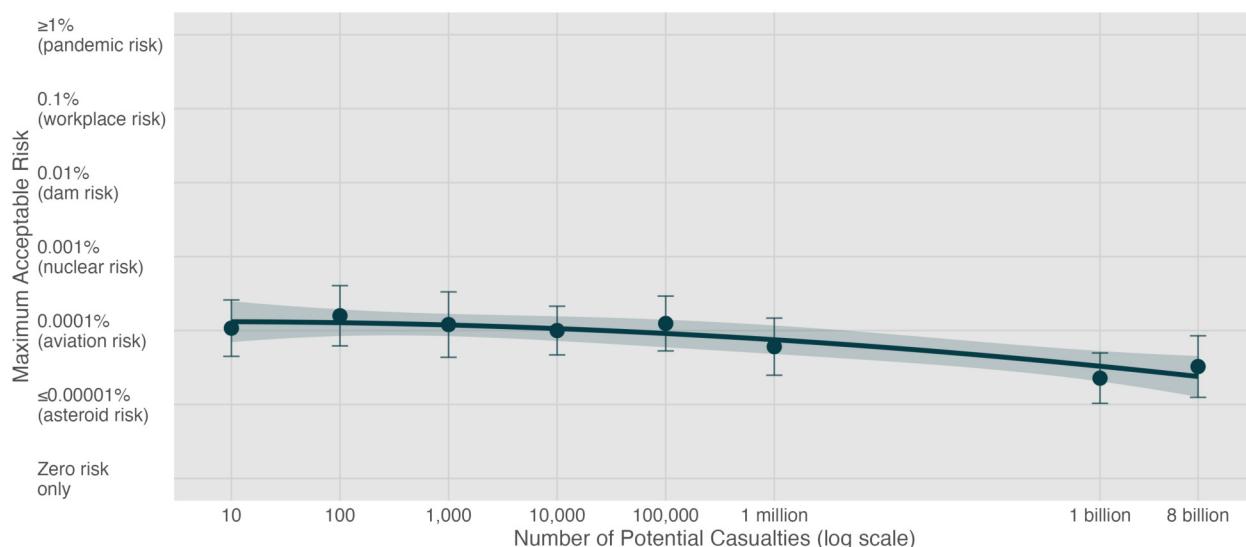
Of 932 participants, 663 (71.1%) passed the consistency check.

Casualty Risk Tolerance (Consistency-Validated)

Risk Tolerance Among Consistency-Validated Respondents (~1 in 1,000,000 / year)

Only participants who passed consistency check (n=663)

What is the HIGHEST probability you would tolerate of AI causing _ casualties each year?



Sample of 663 Australian adults who passed consistency check.

Fieldwork: August 14-30, 2025.

Figure 22: Casualty risk tolerance among participants who passed consistency check

The pattern remains consistent: participants tolerate extremely low risks (~1 in 1,000,000) regardless of casualty scale, even when restricted to those who demonstrated response consistency.

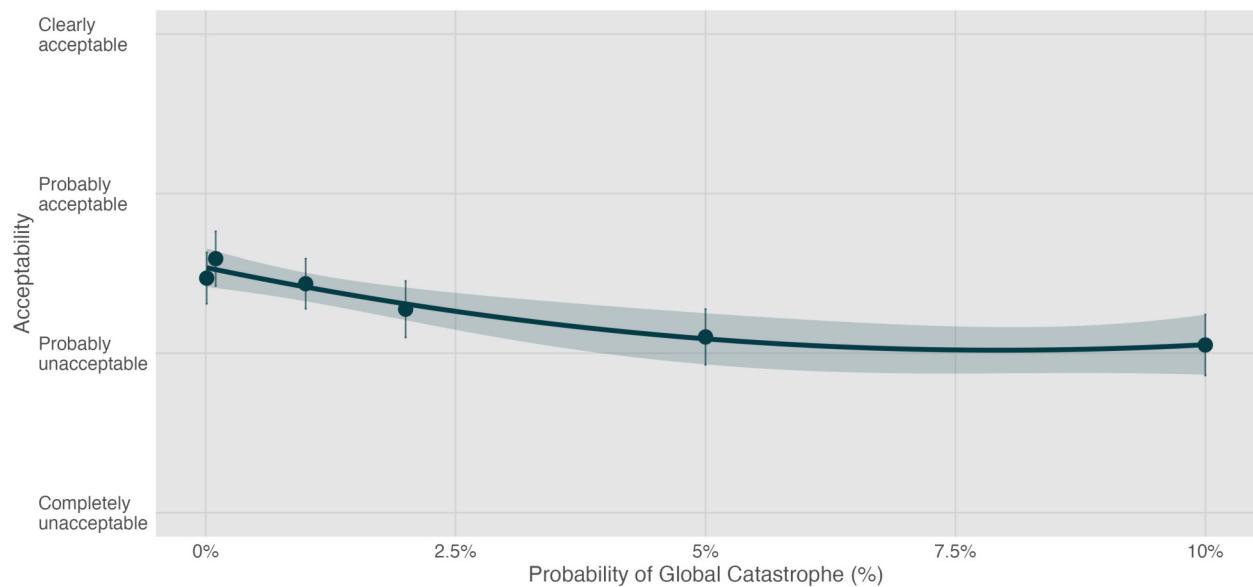
Risk-Benefit Trade-offs (Consistency-Validated)

Risk-Benefit Trade-offs (Consistency-Validated Only)

Only participants who passed consistency check (n=663)

An AI with ___% chance of catastrophe but 60% chance of solving climate change.

This trade-off is:



Sample of 663 Australian adults who passed consistency check.

Fieldwork: August 14-30, 2025.

Figure 23: Risk-benefit trade-offs among consistent responders


Participants who passed the consistency check show similar extreme risk aversion, finding even small probabilities of catastrophe unacceptable despite major benefits.

Temporal Trade-offs (Consistency-Validated)

Temporal Trade-offs (Consistency-Validated Only)

Only participants who passed consistency check (n=663)

If delaying AI by ____ years reduces risk from 5% to 0.5%, this delay would be:



Sample of 663 Australian adults who passed consistency check.
Fieldwork: August 14-30, 2025.

Figure 24: Temporal trade-offs among consistent responders

The willingness to delay AI deployment for safety remains high among consistency-validated respondents.

Statistical Comparison

Table 5: Comparison of risk tolerance measures between full and consistency-validated samples

Measure	Full Sample (n)	Full Sample (Mean ± SE)	Consistent Only (n)	Consistent Only (Mean ± SE)
Casualty Risk Tolerance	722	2.91 ± 0.07	663	2.88 ± 0.07
Risk-Benefit Acceptability	830	2.31 ± 0.03	614	2.32 ± 0.04
Temporal Delay Worthiness	932	4.17 ± 0.03	663	4.21 ± 0.04

Key Finding: The results are robust to excluding participants who failed the consistency check. The extreme risk aversion finding holds even when restricted to participants who demonstrated clear understanding and consistency in their responses.